

Assignment 1

Anhelina Mahdzyar, Boyang Fu, Yao Shi

February 27, 2018

1 Part 1 - Understanding the methods

1.1 Section a

When the agent begins in the maze program, it does not initially know which cells are blocked. It has to survey its surroundings as it takes each proceeding step. In Figure 8, the agent takes a step to the east rather than north because A* would calculate that it would take a shorter distance to the goal to go east, rather than first north and then east. Our agent can only see which cells are blocked at each subsequent step, so at its initial position, it cannot see that the path directly east is blocked both to the right and from above. Only once it moves east will it realize that it is not the correct path, and then move back to go north.

1.2 Section b

In this gridworld, the agent is guaranteed to reach the target of the maze (goal) or discover it is impossible to as A* search leads our agent through every possible cell in the grid to the goal. If the agent detect that the goal state is blocked on all sides by obstacles such as the grid edge or blocked cells, in the next A* search, the open list would be empty before the the algorithm pushing the goal stage node into the open list, and therefore we know we can't reach to the goal stage. Otherwise, the agent will keep moving along each cell until it reaches goal.

Use induction to prove: "The number of moves of the agent until it reaches the target or discovers that this is impossible is bounded from above by the number of unblocked cells squared.

Let n be the total number of cells, let m be the number of unblocked cells, and let k be the total number of moves the agent took

Base case: $n = 2$, when there is only start state and goal state. If the goal state is blocked, then the total move would be $k = 0$, $m = 1$.

$$k < m^2 - - - \text{check}$$

If the goal state is unblocked, then $m = 2$, $k = 1$, because we directly move from start to goal with a single step.

$$k < m^2 - - - \text{check}$$

Inductive step: For $n > 2$, Suppose

$$k < m'^2$$

is true for all $m' \in [1, m)$, need to prove

$$k < m^2$$

is also true. We know that for each unblocked cell, the agent would at most visit it twice (the first time follow the path to the goal, and later come back with the updated blocked information). Therefore, $k \leq 2m$.

$$\left(\frac{2m}{m^2}\right) = \left(\frac{2}{m}\right)$$

, since $m > m' \geq 1$, therefore $m \geq 2$, and further

$$\left(\frac{2}{m}\right) \leq 1$$

. Thus $k/2m \leq 1$. Therefore $k < m^2$ hold true. Based on the base step and inductive step, we know $k < m^2$ for all $m \geq 1$.

2 Part 2 - The Effects of Ties

After implementing tie-breaking in regards to the g-value, we ran our project 10 times for Repeated Forward A* and compared the average run time. In our A* code, implementation of tie-breaking begins around line 254. Here are the run-times:

Repeated Forward A*

Choosing smaller g-value Average run time: 1.733 s

Choosing larger g-value Average run time: 0.584 s

We observed that tie-breaking by the larger g-value proved to be almost three times faster than the other method. From the A* algorithm, $f(s) = g(s) + h(s)$. In terms of the tie-breaking, we can conclude that having a bigger g-value also means having a smaller h-value. In our case, the g-value represents the cost of the path the agent has finished while the h-value is a prediction of the best case in the future which can potentially increase a lot. In other words, having a larger h-value gives the agent more risks of increasing its overall cost in the future. With the same f-values, the one with the larger g-value will more likely to finish its path with less cost to be added on its original prediction. It is also why the method of choosing the larger g-value is faster than the method of choosing the smaller g-value.

3 Part 3 - Forward vs Backward

For this part, we compared Repeated Forward and Backward A* with tie-breaking in favor of cells with larger g-values. We used our results for Forward A* from the previous question. Here are our results:

Choosing larger g-value for tie-breaking and set the graph with 0.2 possibility to be blocked and 0.8 possibility to be unblocked **Repeated Forward A*** Average run time: 0.584 s

Repeated Backward A* Average run time: 0.862 s

We observed that Repeated Forward A* ran faster than Backward A*. **Explanation:**

This is an example of running a maze whose obstacle is closer to the start stage using repeated forward A* algorithm. (**Note: the blue area represents the expanded nodes.)




Figure 1: figure 1.1

This is an example of running the same maze using repeated backward A* algorithm:




Figure 2: figure 1.1

Intuitively from the two graphs, we notice that repeated backward A* search would expand more grid. Formally, suppose during one iteration of A* search, we start the search from a place very close to the already detected vertical obstacle, then A* search would find the break point vertically step by step, in which case each step start over from the beginning to the next vertical point of the obstacle. Therefore, the further the obstacle is from the start point, the more steps it would take to get to the obstacle. As a result, since Repeated Forward A* search start searching from the agent's position, and since the obstacle is always first detected by the agent, the agent are more likely to be closer to the known obstacle at most time. Therefore, comparing to repeated backward A* search, it would generally expand less cell in total to find the path to the goal stage.

4 Part 4 - Heuristics in the Adaptive A*

This part asks us to prove that Manhattan distances are consistent in grid worlds in which the agent can move only in the four main compass directions. The heuristic function h is said to be consistent if

$$\forall(n, a, n') : h(n) \leq c(n, a, n') + h(n'),$$

where $c(n, a, n')$ is the step cost for going from n to n' using action a .

In this case, the step cost for each move is 1, therefore,

$$c(n, a, n') + h(n') = h(n') + 1$$

, and $h(n)$ is either $h(n') + 1$ or $h(n') - 1$, whenever the agent makes the next move.

In the case of $h(n) = h(n') + 1$, $h(n) = c(n, a, n') + h(n')$.

In the case of $h(n) = h(n') - 1$, $h(n) < c(n, a, n') + h(n')$.

Thus, the statement "the Manhattan distances are consistent in grid-worlds in which the agent can move only in the four main compass directions" is true.

The second part asks us to prove that Adaptive A* leaves initially consistent h -values consistent even if action costs can increase. The heuristic function h is said to be consistent if

$$\forall(n, a, n') : h(n) \leq c(n, a, n') + h(n'),$$

where $c(n, a, n')$ is the step cost for going from n to n' using action a .

For each Adaptive process, let the step cost for each move to be denoted by the variable A where A is an integer greater than 0, thus,

$$c(n, a, n') + h(n') = h(n') + A$$

As is known, for Adaptive A* search, $h_{new} = g(s_{goal}) - g(s)$, and we need to prove:

$$h_{new}(s) \leq h_{new}(s') + A,$$

which is equal to:

$$\begin{aligned} g(s_{goal}) - g(s) &\leq g(s_{goal}) - g(s') + A \\ \Rightarrow g(s') &\leq g(s) + A, \end{aligned}$$

which is true.

Thus, the statement "Adaptive A* leaves initially consistent h-values consistent even if action costs can increase" is true.

5 Part 5 – Heuristics in the Adaptive A*

We ran our project 10 times for the Repeated Forward A* and the and Adaptive A*, and compared the average run time. In our A* code, implementation of tie-breaking begins around line 250. Here are the run-times:

Repeated Forward A* Average run time: 1.733 s

Adaptive A* Average run time: 0.584 s

We observed that the repeated forward A* search is much faster than the Adaptive A* search. Both of them break ties among cells with the same f-value in favor of cells with larger g-values and remaining ties in an identical way. Compared between these two algorithms, the Adaptive A* will not be much faster than the Repeated A* because it avoids certain path on the way, while it has to traverse the whole recorded list whenever it updates its next possible path. As for a result, the Adaptive A* only saves a little time from updating its new path but wastes relatively more time traversing list saved whenever it has too, under the condition of each step cost is 1. We also highly suspect that the reason why Adaptive A star search didn't beat repeated Forward A star search is that our obstacle is randomly generated, so it is technically not a "maze", and it might be the reason that the Adaptive A* search didn't show its' advantage in avoiding unnecessary paths. Therefore, the Adaptive A* is slower than the Repeated A* in this case.

6 Part 6 – Memory Issues

In regards to memory consumption, there are a couple of strategies that can be taken in order to reduce it. One method would be to use different variable types in order to save on some space. One simple example, for a variable that indicates whether a cell is blocked or not, this should be a simple boolean value rather than a numerical value. Another would be to limit the number of variables associated with our node class and store less information in general in each node. One of the most important ways is to examine our data structures and see if we may have implemented any other types to save on memory consumption. Another way is to avoid generating the entire map on the agent's side, and only update the nodes that the agent observes. By storing the fog of war, we are unnecessarily using space when the agent can simply create the map as it walks along the grid-world. Our implementation does this but we can see other versions of this project unnecessarily doing this. Furthermore, we can even put the information of all fields inside of one field. For example, our old class cell is written as:

```
class Cell:
    def __init__(self, xPos, yPos, if_blocked, ifVisited=False):
        self.x = xPos
        self.y = yPos
        self.ifBlocked = if_blocked
        self.visited = ifVisited
```

For the sake of saving memory, our new class cell can be written as the following with get's methods:

```
class Cell:
    def __init__(self, xPos, yPos, if_blocked, ifVisited):
        # xPos is the 4-digit x-coordinate, yPos is the 4-digit y-coordinate
        # if_blocked and ifVisited either 1 or 0 in order to indicate a boolean variable
        self.xxxxxyyybv = xPos*1000000+yPos*100+if_blocked*10+ifVisited

    def getx(self):
        temp = self.xxxxxyyybv
        return int(temp / 1000000)

    def gety(self):
        temp = self.xxxxxyyybv - getx(self) * 1000000
        return int(temp / 100)

    def getIfBlocked(self):
        temp = self.xxxxxyyybv - getx(self) * 1000000 - gety(self) * 100
        return int(temp / 10)

    def getIfVisited(self):
        temp = self.xxxxxyyybv
        return temp % 2
```

In this way, each object we create will just be an int type and an int type has a size of 4 bytes and we can do something similar to the class node. Let's take 4MB as $4e6$ bytes, then the maximum number of grids we can have will be $4e6/4/2 = 500000$. $\sqrt{500000} = 707$, which means we will can support a 707x707 grid-world for the worst case. For a 1001x1001 grid-world, it will take $1001*1001*4*2 = 8016008$ bytes of memory which is roughly 8 MB of memory.

*For last question, we have to figure out the lowest amount of memory that one node can hold (examine our current node class and see if there is anything we can trim) and then divide 4MBytes by that number to see the maximum number of cells and thus grid-world that can be generated.